
Patient data storage security

Optinomic GmbH

12 October 2013

The prevalence of mHealth applications is increasing rapidly in the European market. The measurement of therapy-relevant patient parameters and the delivery of therapeutic advice via patients' own mobile devices (smartphones and tablets) poses new challenges in the area of data protection. In this paper we present a proposed plan for patient data management in the Optinomic survey and cognitive testing management framework. This plan is intended to provide strong data protection guarantees while maintaining flexibility and usability for patients and therapists. A cryptographically secure mechanism is used to ensure that patient survey responses can be collected from a patient's own mobile device, but can then only be associated with patient identifying information within the context of a secure server maintained behind a clinic firewall.

Data storage for mHealth

There is good evidence that monitoring of patient mental states and delivery of therapeutic material to patients outside of a clinical setting can lead to improved therapeutic outcomes. The ubiquitous use of mobile devices such as smartphones and tablets provides a clear opportunity to provide such "extra-clinical" support to patients, and modern web application technology permits the development of frameworks that can help therapists to manage the information collected from their patients and to monitor patient states to aid in early intervention in potential problems.

Optinomic is developing a framework to provide adaptive monitoring of patient states, treatment compliance and therapeutic outcomes via surveys and simple cognitive tests delivered to smartphones, tablets or patients' desktop computers. The Optinomic system will integrate with existing Clinical Information Systems, will aid in the collection of legally mandated patient information (e.g. the intake surveys required for the treatment of substance abuse disorders), and will use modern statistical methods to enable therapists and researchers to model patient progress more effectively.

Collection of data using mobile devices and inter-

net technologies has serious implications for compliance with data protection regulations. Patient data is no longer rigidly confined within a clinic's information technology systems where it can be monitored and controlled. This means that careful analysis is needed to ensure that whatever methods are used to collect patient data provide strong data protection guarantees.

Security implications

Within the framework of the applications that we propose, there are two classes of information that must be considered. First, there is *patient identifying information*: patient names, unique identifiers used within Clinic Information Systems, gender, date of birth, treatment history, email addresses, and so on. This information must be protected and should not be accessible to any entities outside of a clinic's security firewall.

The second class of data comprises patient responses to surveys and results of cognitive tests. This data has no *intrinsic* confidentiality requirements – its confidentiality requirements arise solely from the possibility of associating such data with the patient from whom it originated. There should be no way to connect patient survey responses to patient identifying information *except* within the setting of an approved system within the clinic firewall.

The core problem is thus that we would like for patients to be able to access Optinomic applications from their own mobile devices and home computers using the browser that they are familiar with, without needing to take any special security precautions themselves, while at the same time guaranteeing a secure separation between patient identifying information and patient survey responses.

Proposed solution

The architecture that we propose is as follows:

- We use a private *therapy management server* lying within the clinic security firewall along with a publicly accessible *survey server*. The therapy management server, being within the clinic firewall,

will be managed according to the security and data protection policies of the clinic.

- All scheduling of patient survey and test activities is done by the therapy management server.
- The *only* link between a patient and his/her survey/test activity is a randomly generated *survey ticket*.
- There is no correlation between survey tickets for the same patient, or between survey tickets for the same survey or test module, and there is no temporal correlation of tickets, preventing correlation attacks based on ticket allocation.

The operation of this architecture is best illustrated with a detailed view of the primary use case: the scheduling, execution and collection of data for a single activation of a survey module for a single patient.

Use cases

The primary use case we will present here is for a therapist scheduling a patient survey activity and the subsequent presentation of the survey to the patient, followed by data collection. This use case highlights most of the pertinent issues surrounding patient data security for our applications. Some further use cases are considered below.

Primary use case

The numbering of the following steps corresponds to the numbered interactions displayed in Figure 1:

1. The therapist interacts with the Optinomic system from a workstation lying within the clinic's existing IT firewall, where it can be managed in accordance with the data protection policies of the clinic. The therapist uses a web browser to connect to the Optinomic therapy management server which also lies within the clinic firewall. Because both the therapist workstation and the therapy management server lie within the clinic firewall, this connection can be made using a normal HTTP connection. If the IT policies of the clinic permit therapists to access clinic systems from outside the clinic firewall using VPN or other secure connections, this path is also an option for connecting to the Optinomic server.
2. The therapist is then able to access patient details, using either a link to an existing Clinic Information

System or an Optinomic patient details DB. All of these components are contained within the clinic firewall. Unless the integrity of the clinic firewall is breached in some way, there is no possibility of compromise of patient details. Enforcing the integrity of this firewall is beyond the scope of the Optinomic application and is the responsibility of the clinic and their existing information system provider.

3. The therapist selects a therapeutic module (survey or cognitive test) to schedule for a patient by browsing through module definitions stored on a module server (the *Module Market*). No patient data is associated with these module definitions and the module server can thus lie outside the clinic firewall (in many cases, we would recommend that module definitions be accessed from a central Optinomic server so that the latest module versions are always available). Access to the module server is initiated by the therapy management server and uses a normal HTTP connection.
4. The critical step in the data flow is the scheduling of the module for a patient. There needs to be an association made between a unique patient identifier and an instance of the activation of the module. The patient also needs to be informed that the module is available to work on. In order to permit the patient to use their own mobile devices to complete surveys without compromising patient identifying information (name, unique patient ID, email address, and so on) to servers outside the clinic firewall, a random *survey ticket* is generated. This survey ticket is a 32-character string of letters and digits that is generated by the therapy management server for each module activation. The only place where a direct link is made between this random string and the patient ID is an entry created in the Optinomic scheduling database which lies within the clinic firewall and is accessed only by the therapy management server. Once the scheduling database entry has been created, the survey ticket and details of the module to be activated are passed to the *survey server*. Note that no patient identifying information is passed to the survey server, which runs on a machine outside the clinic firewall and can be accessed directly by patients from their own mobile devices or desktop computers. All processing on the survey server is done in reference to these randomly generated survey tickets. Once the survey server has been notified of the module activation, the therapy man-

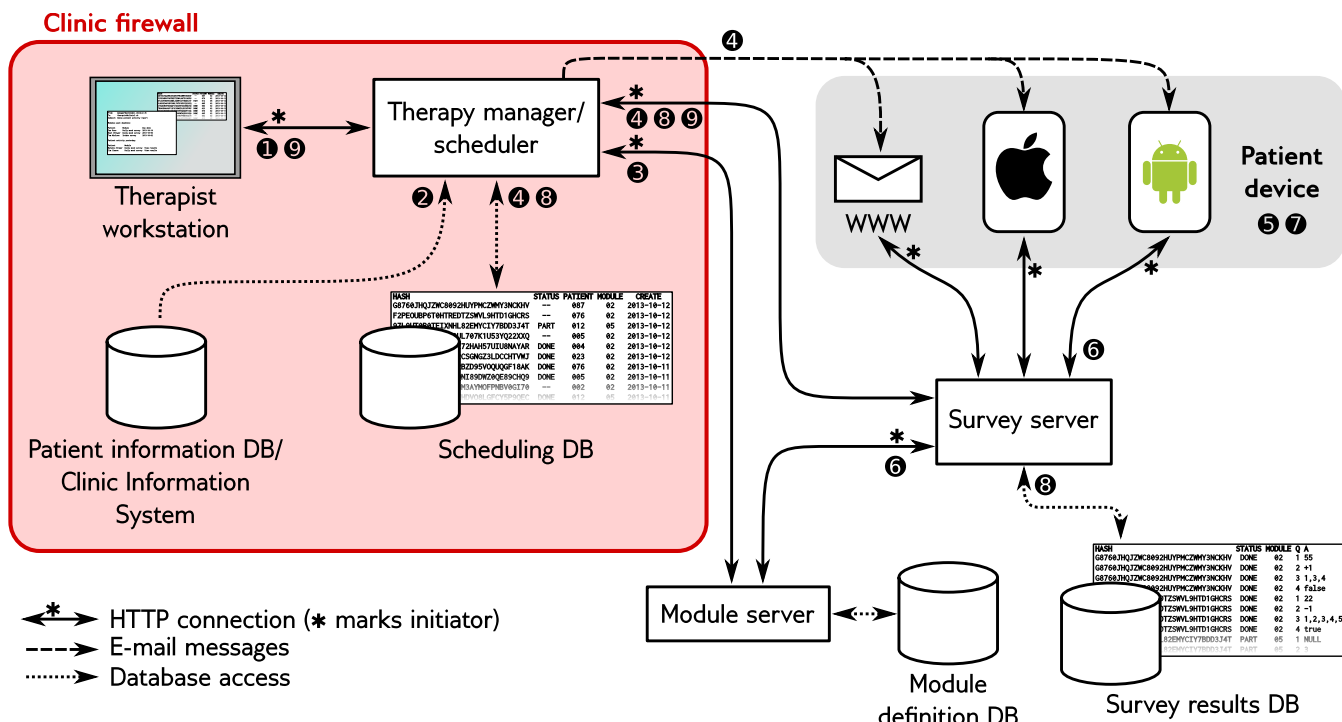


Figure 1: Data flow diagram for primary survey scheduling use case. See text for step-by-step explanation.

agement server sends an email to the patient's registered email address notifying the patient that a new module activation is available. Note again that patient email information is *not* sent to the survey server, although the email sent to the patient does include the survey ticket, in order to provide a link that the patient can follow to complete the survey. For clients who feel that their email infrastructure does not provide adequate levels of security, we suggest using a third party email service, such as Amazon's AWS, that can be accessed via an encrypted HTTPS connection. This measure would protect against attacks that attempt to sniff outgoing emails from the clinic's email server in order to capture patient emails, although it does require the clinic to trust the third party provider.

5. At some point later in time, the patient accesses their email, using whatever mobile device or email client that they normally use and they see that they have a message from the Optinomic therapy manager. This message contains a URL pointing to the relevant survey page, based on the survey ticket generated by the therapy manager (e.g. <http://www.optinomic.ch/G8760JHQJZWC8092HUYPMCZWY3NCKHV>).
6. When the patient follows the URL in the email message, the survey server generates HTML pages and Javascript code to run the relevant survey or cogni-

tive test on the patient's mobile device or computer. Because the survey contents is accessed only via the random survey ticket, this connection can be made via a normal HTTP connection.

7. The patient completes the survey on their mobile device. When they are done, the survey results are submitted to the survey server.
8. When the survey server receives the results from a survey, it saves them to a survey results database. This database is keyed by the random survey ticket, so there is no means by which survey results can be associated with patient identities¹. Once the survey results are saved, they are available to the therapy management server, which polls the survey server on a regular basis (every few minutes or so) to request the survey tickets for any surveys that have been completed recently. (The therapy management server determines this information by polling instead of being notified directly by the survey server in order to avoid the need for the survey server to establish an incoming connection crossing the clinic firewall.) When the therapy management server determines that a survey has been completed, it updates the scheduling database to indicate that the patient has com-

¹The only exception to this statement is (obviously) if the module directly requests patient identifying information as part of a survey.

pleted the assigned activity (at this point, the management server could also send alerts to the therapist, trigger follow-on activities for the patient, and so on).

9. At some point later in time, the therapist once again connects to the therapy management server to view the survey results. The therapy management server is able to access these results from the survey server using a normal HTTP connection. Because the therapy management server has access both to the survey results database (via the survey server) and the scheduling database, it is able to associate survey results to patient identities using the link between the random survey tickets and patient identities stored in the scheduling database.

Other use cases

There are a range of other use cases that are of interest in terms of data security issues. Here we mention three, to indicate how typical uses of the Optinomic system are unlikely to result in compromise of patient data confidentiality.

Survey interruption and resumption For longer surveys, it may be desirable for patients to be able to interrupt and later resume survey completion. There are two possible approaches here: if we permit activities to be paused and resumed only on the same device, then intermediate survey results can simply be stored in a browser cookie for later access. On the other hand, if a patient wants to partially complete a survey on one device (a home computer, perhaps) then finish the survey on another device (a smartphone, for example), then partial survey results can be saved to the survey server. Neither of these approaches risks compromise of the link between patient identity and survey results.

Patient monitoring The survey completion messages sent from the survey server to the therapy management server allow for prompt notification of the therapist of when patients complete certain activities. Because all scheduling and survey completion information is available in the scheduling database, the therapy management server can easily generate reports of patient activity for therapist monitoring, and can also alert the therapist to any concerning patterns of patient non-compliance. All data processing that requires association of survey results with patient identities occurs solely within the clinic firewall, so there is no risk of compromise of patient confidentiality.

Temporal change One of the major benefits of the type of extra-clinical patient monitoring proposed by Optinomic is that it enables ongoing monitoring of patient progress and emotional states. Automated scheduling of the collection of patient data via lightweight daily “mini-surveys” permits the accumulation of long time series of patient data. Combined with simple data visualisation, this allows the therapist to demonstrate therapeutic progress to patients in a very direct way: for example, “In the last month, your daily self-assessment mood scores have been consistently higher than in any of the preceding three months”. Again, for any data analysis of this type, all data processing requiring association of survey results with patient identities occurs solely within the clinic firewall, so there is no risk of compromise of patient confidentiality.

Conclusions

The proposed split between the storage of patient identifying information and survey/test results, using randomly generated survey tickets to provide the link between the two types of data, provides strong protection against breaches in patient data confidentiality (assuming, of course, that the underlying IT security infrastructure is not itself compromised). Within the Optinomic framework, there is no flow of patient identifying data out of the clinic firewall, except for the small possibility of associating patient email addresses and survey tickets if the security of the clinic email system is breached. Apart from this email sniffing attack vector, there is no possibility of making a link between patient survey results and patient identities outside of the firewall, since the only information associating patient identities and survey tickets is stored in the Optinomic scheduling database which is situated within the clinic firewall.